
◆ What This Project Is (Technical but Simple)

We are building an **AI-powered knowledge system** that can understand and answer questions from company documents.

Instead of using traditional keyword search, it uses **semantic understanding** to find the right information and generate accurate answers.

It combines two things:

1. Intelligent search
2. AI-based answer generation

◆ Why We Are Building It

In most companies:

- Documents exist in PDFs, manuals, SOPs, reports
- Searching is manual and time-consuming
- Keyword search fails when wording is different
- Employees repeatedly ask the same questions

We want a system that:

- Understands meaning, not just keywords
- Finds the most relevant document sections
- Generates structured answers
- Provides references

◆ How It Works (Simple Technical Flow)

Step 1 – Document Processing

We upload company documents.

The system:

- Cleans the text
- Breaks it into smaller chunks
- Converts each chunk into numerical embeddings

Step 2 – Vector Search

All document chunks are stored in a vector database.

When a user asks a question:

- The question is also converted into an embedding

- The system finds the most similar document chunks

Step 3 – AI Answer Generation (RAG)

The retrieved document chunks are sent to an AI language model.

The model generates an answer based only on those chunks.

Step 4 – Output

The system returns:

- Answer
- Source references
- Confidence level

This approach is called **Retrieval-Augmented Generation (RAG)**.

◆ Why It Is Different from Normal AI Chatbots

Normal chatbot:

- Can answer from general internet knowledge
- May hallucinate (make up answers)
- No reference validation

Our system:

- Uses only company documents
- Answers are grounded in retrieved data
- References are shown
- If data is missing, it can refuse to answer

This makes it enterprise-safe.

◆ Core Components (Simple Architecture)

The system has 4 main parts:

1. Ingestion Module
Handles document parsing and chunking
2. Embedding & Vector Database
Stores semantic representations
3. RAG Engine (LLM Layer)
Generates context-based answers
4. API Layer
Exposes endpoints for queries and integration

Additionally:

- Evaluation module measures accuracy
 - Monitoring tracks latency and cost
-

◆ **What Problem It Solves Technically**

- Improves search accuracy using semantic similarity
- Reduces hallucination through grounding
- Adds explainability via citations
- Makes document knowledge programmatically accessible

It turns static documents into an intelligent system.

◆ **Future Scope (Technical View)**

Later we can add:

- Hybrid search (keyword + vector)
 - Feedback learning loop
 - Role-based access control
 - Agent workflows (multi-step reasoning)
 - Integration with internal systems
 - Client-facing deployment
-

◆ **One-Line Technical Summary**

“It is a Retrieval-Augmented Generation system that converts enterprise documents into a searchable, grounded AI knowledge platform.”
